

**Recommender Systems with Artificial Intelligence: Architectures, Algorithms,
and Emerging Trends**

Meet Ashokkumar Joshi

Department of Data Science and Artificial Intelligence

Campbellsville University

ABSTRACT

Recommender systems have become fundamental infrastructure in the digital economy, enabling personalized experiences across e-commerce, entertainment, social media, education, and healthcare. The convergence of artificial intelligence (AI) and recommendation technology has catalyzed a paradigm shift from simple heuristic approaches to sophisticated, data-driven models capable of capturing intricate user behaviors and item relationships. This paper presents a comprehensive survey of AI-powered recommender systems, systematically reviewing foundational paradigms including collaborative filtering, content-based filtering, and hybrid methods, while thoroughly examining advanced AI techniques such as matrix factorization, autoencoders, recurrent neural networks, convolutional neural networks, transformer-based architectures, graph neural networks, reinforcement learning, and large language models (LLMs). We analyze the evolution from memory-based to model-based approaches and from static to dynamic, context-aware recommendation. Key application domains, evaluation methodologies, and open challenges—including cold-start problems, data sparsity, scalability, privacy, fairness, and explainability—are examined in depth. This survey also highlights emerging research directions in conversational AI-based recommendation, federated learning for privacy-preserving systems, and cross-domain transfer learning. The findings provide a structured roadmap for researchers and industry practitioners seeking to design and deploy next-generation intelligent recommender systems.

Keywords: Recommender systems, collaborative filtering, content-based filtering, matrix factorization, deep learning, graph neural networks, knowledge graphs, transformers, reinforcement learning, federated learning, large language models, explainability, fairness.

I. INTRODUCTION

The information age has gifted users with access to virtually unlimited digital content, yet this abundance has paradoxically increased the difficulty of discovery. A customer browsing a modern e-commerce platform may face billions of product listings, while a subscriber to a video streaming service encounters millions of titles. Without intelligent mediation, the cognitive burden of navigating such spaces becomes prohibitive. Recommender systems (RS) were conceived to address precisely this challenge—they act as intelligent agents that learn user preferences and proactively surface relevant items, substantially improving user experience and driving platform engagement [8].

The origins of recommendation technology can be traced to the early 1990s, when systems such as Tapestry and GroupLens pioneered collaborative filtering techniques that aggregated user ratings to generate personalized suggestions [1]. Over subsequent decades, the field evolved rapidly, incorporating machine learning models, demographic data, item metadata, and contextual signals to produce progressively more accurate recommendations. The Netflix Prize competition (2006–2009) marked a watershed moment, demonstrating that ensemble methods and latent factor models—particularly matrix factorization—could substantially outperform baseline approaches, spurring widespread academic and industrial interest [2].

The integration of deep learning into recommender systems around 2016 further transformed the landscape. Researchers demonstrated that neural architectures could learn nuanced, non-linear representations of user and item interactions that were far richer than those captured by shallow models [3], [4]. This trend accelerated as attention mechanisms, graph neural networks (GNNs), and transformer architectures were adapted for recommendation tasks, leading to systems that could model sequential behaviors, social relationships, and multi-relational knowledge with unprecedented fidelity [5], [6].

More recently, the emergence of large language models (LLMs) such as GPT-4 and LLaMA has opened entirely new possibilities for recommendation, enabling conversational interfaces, natural language explanations, and zero-shot generalization that were previously intractable [7]. Simultaneously, growing concerns about privacy, fairness, and transparency have driven research into federated learning, differential privacy, debiasing algorithms, and explainable AI for recommender systems.

The scope of RS application has expanded commensurately with their capabilities. Beyond familiar use cases of product recommendation (Amazon), movie suggestion (Netflix), and social content curation (TikTok), AI-driven systems now power personalized healthcare advice, educational content delivery, financial product matching, news aggregation, and scientific literature discovery [8]. In each domain, the fundamental challenge remains consistent: accurately model the preferences, behaviors, and contexts of individual users to surface items of maximum relevance and value.

This paper provides a comprehensive, structured overview of the current state of AI-driven recommender systems. Our contributions include: (1) a systematic review of foundational and advanced recommendation paradigms; (2) an analysis of deep learning, graph-based, and language model techniques applied to recommendation; (3) a discussion of key evaluation frameworks and application domains; and (4) a critical examination of open challenges and future research directions [9]. The remainder of this paper is organized as follows. Section II covers foundational approaches and machine learning techniques. Section III examines deep learning architectures. Section IV discusses graph-based and knowledge-enhanced methods. Section V reviews reinforcement learning and LLM-based approaches. Section VI presents evaluation frameworks and applications. Section VII analyzes open challenges and emerging trends. Section VIII concludes the survey.

II. FOUNDATIONS AND MACHINE LEARNING APPROACHES

A. Collaborative Filtering

Collaborative filtering (CF) remains one of the most studied and widely deployed paradigms in recommendation research. The central hypothesis of CF is that users who have agreed in the past will tend to agree in the future, enabling preference inference across a user community without explicit knowledge of item content [9]. CF methods are broadly divided into memory-based and model-based approaches.

Memory-based CF operates directly on user-item interaction matrices. User-based methods identify users whose historical preferences closely match those of the target user and aggregate their ratings for unseen items [1]. Item-based methods, introduced by Amazon's recommendation team, compute item-item similarities and recommend items similar to those the user has already engaged with [10]. While computationally interpretable, memory-based methods suffer from well-documented scalability limitations and poor performance in sparse interaction matrices, motivating the transition to model-based approaches.

Model-based CF overcomes many of these limitations by learning compact latent representations that capture underlying preference patterns. Matrix factorization (MF) techniques, popularized during the Netflix Prize [2], decompose the user-item interaction matrix into low-dimensional user and item embedding vectors whose dot product approximates the observed interaction. Regularized singular value decomposition (SVD) and alternating least squares (ALS) are canonical MF algorithms. Extensions such as SVD++ incorporate implicit feedback signals, while temporal MF models account for drift in user preferences over time [19].

Factorization machines (FMs) generalized matrix factorization by modeling interactions between arbitrary feature combinations through factorized parameters [11]. This made FMs broadly applicable to recommendation scenarios with rich side information. Field-aware Factorization Machines (FFM)

further improved performance on industrial click-through rate (CTR) prediction tasks, demonstrating the value of capturing field-specific feature interactions at scale.

B. Content-Based Filtering

Content-based filtering (CBF) approaches recommendation as a supervised learning problem: given a user's interaction history, learn a profile that captures preferences for item features, then recommend items whose features closely match this profile [12]. CBF does not require data from other users, making it immune to the cold-start problem for existing users and particularly well-suited to domains with rich item metadata such as news, academic papers, and music.

Classic CBF methods represent items using TF-IDF vectors over textual features and learn user profiles as weighted combinations of engaged item representations. Modern CBF systems employ pretrained language models such as BERT [13] to construct semantically rich item embeddings and adapt transformer encoders to model fine-grained user interest from interaction sequences. Despite their strengths, content-based methods tend to over-specialize on known user preferences, resulting in limited serendipity and reduced capacity to discover items outside the user's established interest space.

C. Hybrid Recommender Systems

Hybrid approaches combine collaborative and content-based signals to mitigate the individual limitations of each paradigm [14]. Hybridization strategies include weighted combinations, feature augmentation, cascade architectures, and unified models that jointly optimize multiple objectives. Industrial systems at companies such as Netflix and Spotify employ highly sophisticated hybrid architectures that integrate dozens of feature sources and model components, achieving substantially better performance than either approach alone. The principled integration of complementary information sources remains a central theme in state-of-the-art recommendation research.

III. DEEP LEARNING ARCHITECTURES

A. Neural Collaborative Filtering

He et al. introduced Neural Collaborative Filtering (NCF), demonstrating that generalized matrix factorization could be naturally expressed within a neural network framework and extended with multi-layer perceptrons (MLPs) to model non-linear user-item interactions [3]. NCF replaced the inner product of traditional MF with a learned interaction function, yielding consistent improvements over MF on standard benchmarks. This work catalyzed a wave of research exploring neural architectures for collaborative filtering, establishing deep learning as the dominant paradigm for recommendation research.

Autoencoders have been applied to CF to learn compact user and item representations from implicit feedback [15]. The AutoRec model encodes item-based or user-based rating vectors through a

bottleneck layer and trains with a reconstruction objective, effectively learning to impute missing interactions. Variational autoencoders (VAE) extended this framework with a probabilistic generative model, improving robustness to noise and enabling principled uncertainty estimation in recommendations [16]. These generative approaches have proven effective for capturing the diverse and often noisy nature of real-world user interaction data.

B. Sequence-Aware and Attention-Based Models

User interaction sequences encode rich temporal dynamics that static CF models ignore. Recurrent neural networks, particularly gated recurrent units (GRU) and long short-term memory (LSTM) networks, were among the first deep architectures applied to model sequential user behavior in recommendation [17]. These models capture the temporal ordering of interactions, enabling prediction of the next item a user is likely to engage with—a task of substantial practical importance for streaming and e-commerce platforms.

Attention mechanisms substantially improved upon recurrent models by allowing the system to weight the influence of different past interactions dynamically. The self-attentive sequential recommendation model SASRec adapted the transformer self-attention mechanism to recommendation, processing the entire interaction history in parallel and capturing long-range dependencies that recurrent models struggle to model efficiently [18]. Subsequent work including BERT4Rec applied bidirectional transformer encoders with masked item prediction, treating sequence modeling analogously to masked language modeling in natural language processing [5]. These transformer-based sequential models represent the current state of the art on sequential recommendation benchmarks and have become a reference architecture for the field.

C. Graph Convolutional Approaches

Graph convolutional networks (GCNs) represent one of the most significant architectural advances in recommendation. By modeling the user-item interaction graph explicitly and propagating information through its structure, GCN-based models capture higher-order connectivity patterns that reflect indirect user similarities and item co-occurrence [6]. Neural Graph Collaborative Filtering (NGCF) demonstrated substantial gains over matrix factorization by incorporating multi-hop graph propagation into the embedding learning process. LightGCN simplified NGCF by removing feature transformation and non-linear activation in each layer, finding that linear neighborhood aggregation already captures sufficient collaborative signals for strong recommendation performance [20].

D. Cross-Feature Interaction Models

Industrial CTR prediction, which underlies most commercial recommender systems, requires modeling complex interactions among sparse categorical features. Deep & Cross Network (DCN) combined a cross network for explicit high-order feature interaction with a deep neural network for implicit interaction modeling [21]. xDeepFM improved upon this approach with a Compressed

Interaction Network (CIN) that learns bounded-degree feature interactions explicitly and combines them with DNN-modeled implicit interactions [22]. These architectures have been widely adopted in production recommendation systems, demonstrating the practical importance of explicitly modeling feature interactions alongside neural representation learning.




Figure 1: NDCG@10 Comparison of Representative AI-Driven Recommendation Models. Values are drawn from original papers [2],[3],[5],[6],[18],[20] on MovieLens-1M and Amazon-Books benchmarks.

IV. GRAPH-BASED AND KNOWLEDGE-ENHANCED METHODS

A. Knowledge Graph Integration

Items in recommendation contexts are rarely isolated entities; they participate in rich webs of relationships representable in knowledge graphs (KGs). KG-enhanced recommender systems leverage these relational structures to improve recommendation quality, especially in cold-start scenarios where interaction data is sparse [23]. Ripple Network propagates user preferences through a knowledge graph, treating user interests as waves that spread along KG connections from seed entities to related concepts. KGCN applies graph convolutional operations on the KG to learn entity embeddings that fuse structural information with recommendation signals.

KGAT jointly models the user-item interaction graph and the KG through an attention-based graph neural network, enabling fine-grained semantic recommendation by distinguishing the importance of different relational paths [24]. Explainability is a notable benefit of KG-enhanced methods: because

recommendation pathways can be traced through the knowledge graph, it becomes possible to provide users with transparent, human-readable justifications for recommendations—a critical property for trust and adoption in high-stakes domains such as healthcare and finance [25].

B. Social Network and Trust-Aware Recommendation

Social influence and trust relationships among users constitute a valuable signal for recommendation. Social recommender systems leverage user-user connections to improve preference estimation, under the assumption that social connections indicate similarity in taste or trustworthiness of opinion [26]. GNN-based models such as GraphRec and DiffNet jointly propagate information over user-user social graphs and user-item interaction graphs, effectively integrating social influence into the recommendation process. These models demonstrate that modeling social dynamics as graph propagation processes substantially improves recommendation accuracy, particularly for users with limited personal interaction history.

V. REINFORCEMENT LEARNING AND LARGE LANGUAGE MODELS

A. Reinforcement Learning for Recommendation

Recommender systems can be naturally framed as sequential decision-making problems, making reinforcement learning (RL) a compelling paradigm for policy optimization [27]. In this formulation, the recommendation agent observes a state representing user context and interaction history, selects an action (the set of items to recommend), and receives a reward signal based on user feedback. Deep Q-Networks (DQN) and policy gradient methods such as actor-critic architectures have been applied to learn recommendation policies that optimize long-term user engagement rather than immediate click-through rate.

DRN modeled news recommendation as an online RL problem, using a dueling bandit framework to explore new items while exploiting known user preferences [28]. RL-based systems are particularly valuable in scenarios where short-term engagement metrics conflict with long-term user satisfaction, such as in avoiding filter bubbles and promoting content diversity. The exploration-exploitation tradeoff inherent to RL naturally enables recommendation systems to balance recommending safe, high-confidence choices with exploring potentially novel and serendipitous items that may better serve long-term user interests.

B. Large Language Models in Recommendation

The advent of large language models has introduced a transformative new paradigm for recommender systems. LLMs possess broad world knowledge, natural language understanding, and in-context learning capabilities that enable reasoning about user preferences from text descriptions alone, without requiring large amounts of task-specific training data [7]. Several approaches have emerged for integrating LLMs into recommendation pipelines. Prompt-based methods formulate recommendation

as a text-to-text task, providing the LLM with user history and asking it to generate item recommendations directly.

P5 unified diverse recommendation tasks—rating prediction, sequential recommendation, explanation generation, and review summarization—within a single text-to-text T5 framework, demonstrating the remarkable versatility of pretrained language models for recommendation [29]. LLM-based recommendation faces unique challenges including hallucination of item identifiers, inference latency, and the difficulty of staying updated with rapidly changing item catalogs. Retrieval-augmented generation (RAG) architectures address some of these limitations by grounding LLM generation in real-time retrieval from item indices, offering a promising direction for scalable and up-to-date LLM-powered recommendation.

C. Conversational Recommender Systems

Conversational recommender systems (CRS) engage users in multi-turn dialogues to elicit preferences and progressively refine recommendations [30]. Early CRS systems relied on slot-filling templates and decision trees; modern systems leverage neural dialogue models and LLMs to support open-domain natural language interaction. The CRS task is typically decomposed into preference elicitation (asking clarifying questions), recommendation generation, and explanation provision. Models combining dialogue modeling with knowledge graph-based recommendation have demonstrated significant improvements over single-component approaches, while GPT-based systems exhibit strong zero-shot generalization to new recommendation domains.

VI. EVALUATION FRAMEWORKS AND APPLICATIONS

A. Evaluation Metrics

Rigorous evaluation is fundamental to research progress in recommender systems. Standard offline evaluation metrics include precision, recall, F1-score, normalized discounted cumulative gain (NDCG), mean average precision (MAP), and hit rate (HR), computed over ranked recommendation lists against held-out interaction data [8]. For rating prediction tasks, mean absolute error (MAE) and root mean squared error (RMSE) are standard metrics. Beyond accuracy, evaluation frameworks increasingly incorporate diversity, novelty, serendipity, coverage, and fairness metrics to assess the broader quality of recommendation [31]. A critical limitation of offline evaluation is that it does not capture the dynamic, interactive nature of real recommendation environments. Online A/B testing remains the gold standard for evaluating systems in production, though it is expensive. Counterfactual offline evaluation and interleaving methods have emerged as practical intermediate alternatives.

Table I: Comparison of Representative AI-Driven Recommendation Models

Model	Year	Core Technique	Task
MF / SVD++	2009	Matrix Factorization	Rating Prediction
FM / FFM	2010	Feature Interaction	CTR Prediction
AutoRec / MultVAE	2015/18	Autoencoder / VAE	Top-N Rec.
NCF	2017	MLP + GMF	Top-N Rec.
YouTube DNN	2016	Deep Neural Network	Video Rec.
GRU4Rec	2016	Recurrent (GRU)	Sequential Rec.
Ripple Net / KGAT	2018/19	KG Propagation	Top-N Rec.
NGCF / LightGCN	2019/20	Graph Convolution	CF Rec.
SASRec	2018	Transformer (Unidirectional)	Sequential Rec.
BERT4Rec	2019	Transformer (Bidirectional)	Sequential Rec.
DCN / xDeepFM	2017/18	Cross Feature Network	CTR Prediction

Model	Year	Core Technique	Task
DRN	2018	Deep RL (Dueling Bandit)	News Rec.
P5	2022	T5 LLM (Pretrained)	Multi-Task Rec.

Abbreviations: Rec. = Recommendation; CTR = Click-Through Rate; CF = Collaborative Filtering; KG = Knowledge Graph; MLP = Multi-Layer Perceptron; GMF = Generalized Matrix Factorization.

B. Application Domains

The application landscape for AI-driven recommender systems spans virtually every domain of digital commerce and content. E-commerce recommendation, as exemplified by Amazon's item-to-item CF system [10], drives significant revenue through cross-selling and upselling recommendations. Streaming service recommendation at Netflix and YouTube is the primary mechanism through which users discover new content, directly influencing churn and engagement metrics [4]. Social media feed ranking on platforms such as TikTok and Instagram employs highly sophisticated AI models to curate personalized timelines that maximize engagement.

In healthcare, recommender systems assist clinicians in identifying relevant literature, suggest personalized treatment protocols, and recommend lifestyle interventions based on patient data. Educational recommender systems deliver personalized learning pathways, adapting content sequencing to individual student performance and learning style [31]. Financial recommender systems suggest investment products, insurance plans, and credit options, subject to stringent regulatory requirements for fairness and explainability. The breadth of these application domains underscores the practical importance of continued advances in AI-driven recommendation research.




Figure 2: Approximate Hit Rate@10 Milestones of Key Recommender System Models (2009–2023). Values are illustrative, compiled from benchmark results reported in cited papers [2],[3],[5],[6],[7],[11],[15],[18],[19],[20],[29] across comparable datasets.

VII. CHALLENGES AND EMERGING RESEARCH DIRECTIONS

A. Cold-Start Problem

The cold-start problem—the difficulty of generating meaningful recommendations for new users or items with no interaction history—remains a fundamental challenge in recommender systems research [9]. For new users, content-based onboarding questionnaires and transfer learning from demographic or behavioral signals offer partial mitigation. For new items, knowledge graph embeddings and zero-shot LLM-based inference can bootstrap item representations without requiring historical interactions [7]. Meta-learning approaches such as MAML (Model-Agnostic Meta-Learning) adapt quickly to new users or items from minimal interaction data, demonstrating promising cold-start performance on standard benchmarks. The cold-start problem is increasingly viewed as a specific instance of the broader challenge of learning generalizable preference representations that can transfer across contexts.

B. Data Sparsity and Scalability

Real-world interaction matrices are extremely sparse; the majority of users interact with only a tiny fraction of available items. This sparsity degrades collaborative filtering performance and amplifies noise in similarity computations. Techniques such as data augmentation through negative sampling, graph augmentation, and dropout regularization partially address sparsity. Scalability is an orthogonal

challenge: as item catalogs grow to billions and user bases reach hundreds of millions, approximate nearest neighbor search, product quantization, and two-stage retrieval-reranking architectures are necessary to maintain low inference latency at production scale [32]. The two-stage paradigm, in which a lightweight retrieval model generates a candidate set subsequently reranked by a more complex model, is now standard in industrial recommendation pipelines.

C. Fairness, Bias, and Explainability

AI-driven recommender systems can perpetuate and amplify biases present in historical interaction data, leading to unfair outcomes for users and items belonging to minority groups [33]. Popularity bias causes systems to systematically over-recommend popular items at the expense of long-tail items. Demographic bias may cause systematic differences in recommendation quality across user groups defined by age, gender, or location. Debiasing methods include inverse propensity scoring, adversarial debiasing, and constrained optimization formulations that incorporate fairness constraints directly into the training objective.

Explainability is essential for user trust, regulatory compliance, and debugging. Attention visualization, path-based reasoning over knowledge graphs, counterfactual explanation methods, and natural language explanations generated by LLMs are active areas of research [25]. Emerging regulations such as the EU AI Act are expected to impose explicit requirements for explainability in high-stakes recommendation contexts, making this research direction increasingly important for practical deployment.

D. Privacy Preservation and Federated Learning

Recommendation systems require access to sensitive user behavioral data, creating significant privacy risks. Federated learning (FL) enables model training across decentralized user devices without centralizing raw interaction data, with only model gradients or compressed updates being communicated to a central server [34]. Differential privacy techniques add calibrated noise to training updates to provide formal mathematical privacy guarantees. Secure multi-party computation protocols enable collaborative model training among competing parties without revealing private data. These privacy-preserving techniques introduce accuracy-privacy tradeoffs that remain an active research frontier, with recent work demonstrating that carefully designed FL protocols can approach the performance of centralized training.

E. Cross-Domain and Transfer Learning

Users interact with multiple platforms and domains whose recommendation systems operate in isolation, forfeiting valuable cross-domain preference signals. Cross-domain recommendation transfers knowledge learned in a data-rich source domain to a data-sparse target domain, mitigating cold-start problems and improving generalization [35]. Transfer learning approaches based on pretrained language models and multi-task learning frameworks have shown substantial promise.

LLMs are particularly well-suited as foundational representations for cross-domain recommendation, given their broad world knowledge acquired from diverse pretraining corpora [7].

F. Session-Based Recommendation

Session-based recommendation addresses the scenario in which user identity is unknown and the only available context is the sequence of items interacted with in the current browsing session [36]. This is a practically important scenario for e-commerce and web platforms where the majority of traffic is anonymous. GNN-based models such as SR-GNN model intra-session item transitions as directed graphs, capturing complex session dynamics that sequential models miss. Self-supervised contrastive learning methods have further advanced state-of-the-art performance on session-based benchmarks, exploiting data augmentation to learn robust session representations without additional supervision. The session-based setting also connects naturally to conversational recommendation, where each dialogue turn can be viewed as a session context update.

VIII. CONCLUSION

This paper has presented a comprehensive survey of AI-driven recommender systems, tracing their evolution from early collaborative filtering methods to contemporary deep learning, graph neural network, reinforcement learning, and large language model paradigms. The field has seen remarkable progress in personalized recommendation accuracy, driven by increasingly sophisticated architectures and the availability of large-scale interaction data.

Key insights from our survey include: (1) deep learning models have consistently outperformed shallow methods by learning rich non-linear user-item representations [3]; (2) graph neural networks have proven highly effective for capturing high-order collaborative signals in user-item and knowledge graphs [6], [20]; (3) transformer and LLM-based approaches are rapidly reshaping the landscape, enabling conversational recommendation and cross-domain generalization [5], [7], [29]; and (4) fairness, privacy, explainability, and scalability remain critical open challenges that the community is actively addressing [25], [33], [34].

Looking forward, the integration of generative AI, multimodal foundation models, and real-time adaptation mechanisms promises to further elevate recommendation quality and breadth. The responsible deployment of these systems—ensuring they serve diverse user populations equitably, respect user privacy, and operate transparently—will require sustained interdisciplinary effort spanning machine learning, human-computer interaction, economics, and policy [8], [9]. The continued evolution of recommender systems will remain central to the design of user-centric digital experiences in the coming decade, with AI serving as the enabling technology for ever-more-personalized and beneficial recommendations.

REFERENCES

- [1] J. A. Konstan, B. N. Miller, D. Maltz, J. L. Herlocker, L. R. Gordon, and J. Riedl, "GroupLens: Applying collaborative filtering to Usenet news," *Commun. ACM*, vol. 40, no. 3, pp. 77–87, Mar. 1997. doi: <https://doi.org/10.1145/245108.245126>
- [2] Y. Koren, R. Bell, and C. Volinsky, "Matrix factorization techniques for recommender systems," *Computer*, vol. 42, no. 8, pp. 30–37, Aug. 2009. doi: <https://doi.org/10.1109/MC.2009.263>
- [3] X. He, L. Liao, H. Zhang, L. Nie, X. Hu, and T.-S. Chua, "Neural collaborative filtering," in *Proc. 26th Int. Conf. World Wide Web (WWW)*, Perth, Australia, Apr. 2017, pp. 173–182. doi: <https://doi.org/10.1145/3038912.3052569>
- [4] P. Covington, J. Adams, and E. Sargin, "Deep neural networks for YouTube recommendations," in *Proc. 10th ACM Conf. Recommender Systems (RecSys)*, Boston, MA, USA, Sep. 2016, pp. 191–198. doi: <https://doi.org/10.1145/2959100.2959190>
- [5] F. Sun, J. Liu, J. Wu, C. Pei, X. Lin, W. Ou, and P. Jiang, "BERT4Rec: Sequential recommendation with bidirectional encoder representations from transformer," in *Proc. 28th ACM Int. Conf. Information and Knowledge Management (CIKM)*, Beijing, China, Nov. 2019, pp. 1441–1450. doi: <https://doi.org/10.1145/3357384.3357895>
- [6] X. Wang, X. He, M. Wang, F. Feng, and T.-S. Chua, "Neural graph collaborative filtering," in *Proc. 42nd Int. ACM SIGIR Conf. Research and Development in Information Retrieval*, Paris, France, Jul. 2019, pp. 165–174. doi: <https://doi.org/10.1145/3331184.3331267>
- [7] Z. Zhao, W. Fan, J. Li, Y. Liu, X. Mei, Y. Wang, Z. Wen, F. Wang, X. Zhao, J. Tang, and Q. Li, "Recommender systems in the era of large language models (LLMs)," *IEEE Trans. Knowl. Data Eng.*, vol. 36, no. 11, pp. 6889–6907, Nov. 2024. doi: <https://doi.org/10.1109/TKDE.2024.3392335>
- [8] G. Adomavicius and A. Tuzhilin, "Toward the next generation of recommender systems: A survey of the state-of-the-art and possible extensions," *IEEE Trans. Knowl. Data Eng.*, vol. 17, no. 6, pp. 734–749, Jun. 2005. doi: <https://doi.org/10.1109/TKDE.2005.99>
- [9] F. Ricci, L. Rokach, and B. Shapira, Eds., *Recommender Systems Handbook*, 2nd ed. New York, NY, USA: Springer, 2015.
- [10] G. Linden, B. Smith, and J. York, "Amazon.com recommendations: Item-to-item collaborative filtering," *IEEE Internet Comput.*, vol. 7, no. 1, pp. 76–80, Jan.–Feb. 2003. doi: <https://doi.org/10.1109/MIC.2003.1167344>
- [11] S. Rendle, "Factorization machines," in *Proc. IEEE Int. Conf. Data Mining (ICDM)*, Sydney, Australia, Dec. 2010, pp. 995–1000. doi: <https://doi.org/10.1109/ICDM.2010.127>
- [12] M. J. Pazzani and D. Billsus, "Content-based recommendation systems," in *The Adaptive Web*, P. Brusilovsky, A. Kobsa, and W. Nejdl, Eds. Berlin, Germany: Springer, 2007, pp. 325–341. doi: https://doi.org/10.1007/978-3-540-72079-9_10

- [13] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "BERT: Pre-training of deep bidirectional transformers for language understanding," in Proc. 2019 Conf. North American Chapter Assoc. Computational Linguistics: Human Language Technologies (NAACL-HLT), Minneapolis, MN, USA, Jun. 2019, vol. 1, pp. 4171–4186. doi: <https://doi.org/10.18653/v1/N19-1423>
- [14] R. Burke, "Hybrid recommender systems: Survey and experiments," User Model. User-Adapt. Interact., vol. 12, no. 4, pp. 331–370, Nov. 2002. doi: <https://doi.org/10.1023/A:1021240730564>
- [15] S. Sedhain, A. K. Menon, S. Sanner, and L. Xie, "AutoRec: Autoencoders meet collaborative filtering," in Proc. 24th Int. Conf. World Wide Web (WWW) Companion, Florence, Italy, May 2015, pp. 111–112. doi: <https://doi.org/10.1145/2740908.2742726>
- [16] D. Liang, R. G. Krishnan, M. D. Hoffman, and T. Jebara, "Variational autoencoders for collaborative filtering," in Proc. 27th Int. Conf. World Wide Web (WWW), Lyon, France, Apr. 2018, pp. 689–698. doi: <https://doi.org/10.1145/3178876.3186150>
- [17] B. Hidasi, A. Karatzoglou, L. Baltrunas, and D. Tikk, "Session-based recommendations with recurrent neural networks," in Proc. 4th Int. Conf. Learning Representations (ICLR), San Juan, Puerto Rico, May 2016. [Online]. Available: <https://arxiv.org/abs/1511.06939>
- [18] W.-C. Kang and J. McAuley, "Self-attentive sequential recommendation," in Proc. IEEE Int. Conf. Data Mining (ICDM), Singapore, Nov. 2018, pp. 197–206. doi: <https://doi.org/10.1109/ICDM.2018.00035>
- [19] Y. Koren, "Factorization meets the neighborhood: A multifaceted collaborative filtering model," in Proc. 14th ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining, Las Vegas, NV, USA, Aug. 2008, pp. 426–434. doi: <https://doi.org/10.1145/1401890.1401944>
- [20] X. He, K. Deng, X. Wang, Y. Li, Y. Zhang, and M. Wang, "LightGCN: Simplifying and powering graph convolution network for recommendation," in Proc. 43rd Int. ACM SIGIR Conf. Research and Development in Information Retrieval, Virtual Event, China, Jul. 2020, pp. 639–648. doi: <https://doi.org/10.1145/3397271.3401063>
- [21] R. Wang, B. Fu, G. Fu, and M. Wang, "Deep & cross network for ad click predictions," in Proc. ADKDD'17: Workshop on Advertising, Data Mining and Knowledge Discovery, Halifax, NS, Canada, Aug. 2017, pp. 1–7. doi: <https://doi.org/10.1145/3124749.3124754>
- [22] J. Lian, X. Zhou, F. Zhang, Z. Chen, X. Xie, and G. Sun, "xDeepFM: Combining explicit and implicit feature interactions for recommender systems," in Proc. 24th ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining, London, UK, Aug. 2018, pp. 1754–1763. doi: <https://doi.org/10.1145/3219819.3220023>
- [23] H. Wang, F. Zhang, J. Wang, M. Zhao, W. Li, X. Xie, and M. Guo, "Ripple Network: Propagating user preferences on the knowledge graph for recommender systems," in Proc. 27th ACM Int. Conf. Information and Knowledge Management (CIKM), Turin, Italy, Oct. 2018, pp. 417–426. doi: <https://doi.org/10.1145/3269206.3271739>

- [24] X. Wang, X. He, Y. Cao, M. Liu, and T.-S. Chua, "KGAT: Knowledge graph attention network for recommendation," in Proc. 25th ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining, Anchorage, AK, USA, Aug. 2019, pp. 950–958. [doi: https://doi.org/10.1145/3292500.3330989](https://doi.org/10.1145/3292500.3330989)
- [25] X. Wang, D. Wang, C. Xu, X. He, Y. Cao, and T.-S. Chua, "Explainable reasoning over knowledge graphs for recommendation," in Proc. 33rd AAAI Conf. Artificial Intelligence, Honolulu, HI, USA, Jan. 2019, vol. 33, pp. 5329–5336. [doi: https://doi.org/10.1609/aaai.v33i01.33015329](https://doi.org/10.1609/aaai.v33i01.33015329)
- [26] W. Fan, Y. Ma, Q. Li, Y. He, E. Zhao, J. Tang, and D. Yin, "Graph neural networks for social recommendation," in Proc. 28th Int. Conf. World Wide Web (WWW), San Francisco, CA, USA, May 2019, pp. 417–426. [doi: https://doi.org/10.1145/3308558.3313488](https://doi.org/10.1145/3308558.3313488)
- [27] X. Zhao, L. Zhang, Z. Ding, L. Xia, J. Tang, and D. Yin, "Recommendations with negative feedback via pairwise deep reinforcement learning," in Proc. 24th ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining, London, UK, Aug. 2018, pp. 1040–1048. [doi: https://doi.org/10.1145/3219819.3219886](https://doi.org/10.1145/3219819.3219886)
- [28] G. Zheng, F. Zhang, Z. Zheng, Y. Xiang, N. J. Yuan, X. Xie, and Z. Li, "DRN: A deep reinforcement learning framework for news recommendation," in Proc. 27th Int. Conf. World Wide Web (WWW), Lyon, France, Apr. 2018, pp. 167–176. [doi: https://doi.org/10.1145/3178876.3185994](https://doi.org/10.1145/3178876.3185994)
- [29] S. Geng, S. Liu, Z. Fu, Y. Ge, and Y. Zhang, "Recommendation as language processing (RLP): A unified pretrain, personalized prompt & predict paradigm (P5)," in Proc. 16th ACM Conf. Recommender Systems (RecSys), Seattle, WA, USA, Sep. 2022, pp. 299–315. [doi: https://doi.org/10.1145/3523227.3546767](https://doi.org/10.1145/3523227.3546767)
- [30] Y. Sun and Y. Zhang, "Conversational recommender system," in Proc. 41st Int. ACM SIGIR Conf. Research and Development in Information Retrieval, Ann Arbor, MI, USA, Jul. 2018, pp. 235–244. [doi: https://doi.org/10.1145/3209978.3210002](https://doi.org/10.1145/3209978.3210002)
- [31] S. Zhang, L. Yao, A. Sun, and Y. Tay, "Deep learning based recommender system: A survey and new perspectives," ACM Comput. Surv., vol. 52, no. 1, pp. 1–38, Feb. 2019. [doi: https://doi.org/10.1145/3285029](https://doi.org/10.1145/3285029)
- [32] J. Johnson, M. Douze, and H. Jégou, "Billion-scale similarity search with GPUs," IEEE Trans. Big Data, vol. 7, no. 3, pp. 535–547, Jul. 2021. [doi: https://doi.org/10.1109/TBDATA.2019.2921572](https://doi.org/10.1109/TBDATA.2019.2921572)
- [33] M. D. Ekstrand, R. Burke, and F. Diaz, "Fairness and discrimination in retrieval and recommendation," in Proc. 42nd Int. ACM SIGIR Conf. Research and Development in Information Retrieval, Paris, France, Jul. 2019, pp. 1403–1404. [doi: https://doi.org/10.1145/3331184.3331380](https://doi.org/10.1145/3331184.3331380)

- [34] Q. Yang, Y. Liu, T. Chen, and Y. Tong, "Federated machine learning: Concept and applications," *ACM Trans. Intell. Syst. Technol.*, vol. 10, no. 2, pp. 1–19, Jan. 2019. doi: <https://doi.org/10.1145/3298981>
- [35] T. Man, H. Shen, X. Jin, and X. Cheng, "Cross-domain recommendation: An embedding and mapping approach," in *Proc. 26th Int. Joint Conf. Artificial Intelligence (IJCAI)*, Melbourne, Australia, Aug. 2017, pp. 2464–2470. doi: <https://doi.org/10.24963/ijcai.2017/343>
- [36] S. Wu, Y. Tang, Y. Zhu, L. Wang, X. Xie, and T. Tan, "Session-based recommendation with graph neural networks," in *Proc. 33rd AAAI Conf. Artificial Intelligence*, Honolulu, HI, USA, Jan. 2019, vol. 33, pp. 346–353. doi: <https://doi.org/10.1609/aaai.v33i01.3301346>